Code Generation

Multiple IRs in the Compiler

Modern Compiler Framework (Clang/LLVM)

- Remember this diagram from our first day?

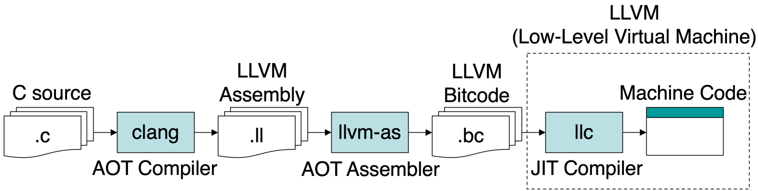

- LLVM Bitcode is in LLVM IR (Intermediate Representation)

Modern Compiler Framework (Clang/LLVM)

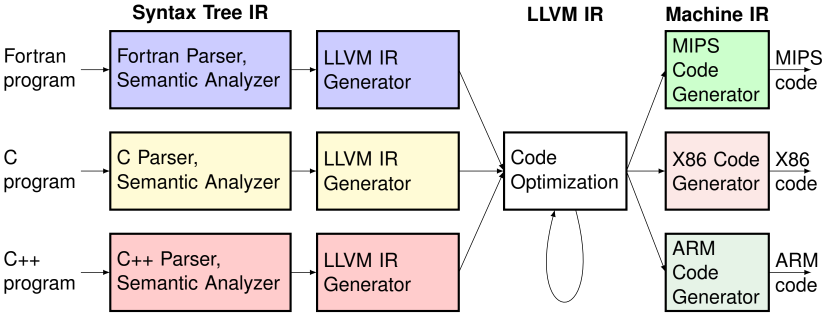

- Common LLVM IR for all languages and backends means:
  - Code optimizations need to be written only once
  - Implementation complexity if O(M + N) instead of O(M * N) (where M = number of frontends, N = number of backends)

Why O(M * N) when no common IR?

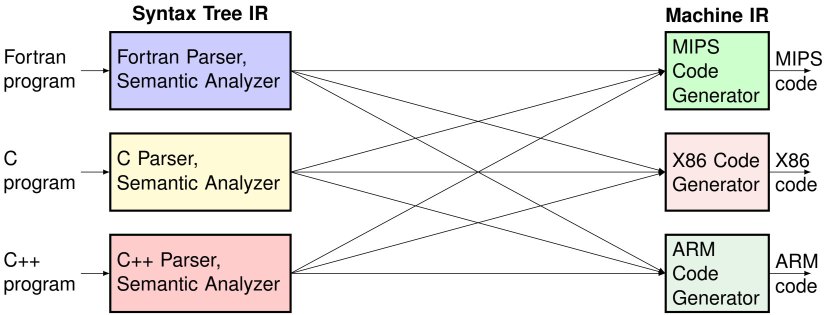

- Must translate M languages to N machine codes
  - Must also do optimizations during each of these translations

High-Level IRs

- Goal: Express the syntax and semantics of source code
- Examples: Abstract Syntax Tree, Parse Tree
- Differs on: Source code programming language
- Uses:
  - Generated by syntax analysis
  - Used by semantic analysis for binding and type checking
  - Language-specific optimizations

High-Level IRs: Language-specific Optimizations

- IR expresses language constructs specific to langauge
  $\rightarrow$ Suitable for language-specific optimizations

High-Level IRs: Language-specific Optimizations

- IR expresses language constructs specific to langauge
  $\rightarrow$ Suitable for language-specific optimizations
- E.g. Devirtualization for languages with polymorphism
  - Polymorphic method calls are indirect jumps using a vtable (Vtable: a table of function pointers specific to each class)
  - **Devirtualization**: changing polymorphic calls to direct calls
  - If runtime object type can be proven (using type inference)
    $\rightarrow$ Indirect call can be changed to direct call
    $\rightarrow$ Direct call can sometimes be inlined into caller method

High-Level IRs: Language-specific Optimizations

- IR expresses language constructs specific to langauge
  → Suitable for language-specific optimizations
- E.g. Devirtualization for languages with polymorphism
  - Polymorphic method calls are indirect jumps using a vtable (Vtable: a table of function pointers specific to each class)
  - **Devirtualization**: changing polymorphic calls to direct calls
  - If runtime object type can be proven (using type inference)
    → Indirect call can be changed to direct call
    → Direct call can sometimes be inlined into caller method
- E.g. Type specialization for languages with dynamic typing
  - Naively, all ops on variables needs dynamic type checking (A big switch/case statement with code for each type)
  - **Type specialization**: generating code for just one type
    - By proving runtime type of variable (using type inference)
    - By profiling type at runtime and generating code for that type

Low-Level IRs

- Goal: Express code in the ISA of an abstract machine
- Examples: Three address code, Static Single Assignment
- Differs on: Language and back-end machine agnostic
- Uses:
  - A common IR that connects front-ends and back-ends
  - Language / machine independent optimizations

Low-Level IRs: Universal Optimizations

- IR is language and machine agnostic
  $\rightarrow$ Do optimizations common to all languages/machines

Low-Level IRs: Universal Optimizations

- IR is language and machine agnostic
  $\rightarrow$ Do optimizations common to all languages/machines

- Deleting code (**dead code elimination**)
  A=2; A=y; $\equiv$ A=y;

Low-Level IRs: Universal Optimizations

- IR is language and machine agnostic
  $\rightarrow$ Do optimizations common to all languages/machines

- Deleting code (**dead code elimination**)
  A=2; A=y; $\equiv$ A=y;
- Deleting code (**common subexpression elimination**)
  A=x+y; B=x+y+z; $\equiv$ A=x+y; B=A+z

Low-Level IRs: Universal Optimizations

- IR is language and machine agnostic
  $\rightarrow$ Do optimizations common to all languages/machines

- Deleting code (**dead code elimination**)
  A=2; A=y; $\equiv$ A=y;
- Deleting code (**common subexpression elimination**)
  A=x+y; B=x+y+z; $\equiv$ A=x+y; B=A+z
- Moving code (**loop invariant code motion**)
  for (i = 0; i < 100; i++) { sum += i + x * y; }
  $\equiv$
  t = x * y;
  for (i = 0; i < 100; i++) { sum += i + t; }

Machine IRs

- Goal: Generate code in the ISA of back-end machine
- Examples: x86 IR, ARM IR, MIPS IR
- Differs on: Back-end machine ISA
- Uses:
  - Register allocation / machine code generation
  - Machine-specific optimizations
    - Strength reduction (replacing op with cheaper op)
    - Vectorization (using CPU vector units if available)
    - ...

Machine IRs: Machine-specific Optimizations

- IR expresses instructions specific to machine
  $\rightarrow$ Suitable for machine-specific optimizations

Machine IRs: Machine-specific Optimizations

- IR expresses instructions specific to machine
  $\rightarrow$ Suitable for machine-specific optimizations

- **Strength reduction**: replacing op with cheaper op
  A=2*a; $\equiv$ A=a«1;

Machine IRs: Machine-specific Optimizations

- IR expresses instructions specific to machine
  $\rightarrow$ Suitable for machine-specific optimizations

- **Strength reduction**: replacing op with cheaper op
  A=2*a; $\equiv$ A=a«1;
- **Vectorization**: using CPU vector units to parallelize
  for (i = 0; i < 64; i++) { C[i] = A[i] + B[i]; }
  $\equiv$
  vec_add_64 vec_register, A, B
  vec_store_64 C, vec_register

Low-Level IRs

Three Address Code

Generic form is **X = Y op Z**

where X, Y, Z can be variables, constants, or compiler-generated temporaries holding intermediate values

- Characteristics
  - Assembly code for an 'abstract machine'
  - Long expressions are converted to multiple instructions
  - Control flow statements are converted to jumps
  - Machine independent
    - Operations are generic (not tailored to specific machine)
    - Function calls represented as generic call nodes
    - Uses **symbolic names** rather than **register names**
      (Actual locations of symbols are yet to be determined)
- Why this form?
  - Allows IR to be generated in a machine-agnostic way
  - Optimizations on IR can be done much more easily
    (Optimizations don't worry about syntactic structure)

Example

- An example:

  x * y + z / w

  is translated to

  t1 = x * y ; t1, t2, t3 are temporary variables

  t2 = z / w

  t3 = t1 + t2

  - Sequential translation of an AST
  - Internal nodes in AST are translated to temporary variables
  - Can be generated through a depth-first traversal of AST

Common Three-Address Statements (I)

- Binary operation statement:

  x = y op z

  where op is an arithmetic or logical operation (binary operation)
- Unary operation statement:

  x = op y

  where op is an unary operation such as -, not, shift)
- Copy statement:

  x = y
- Unconditional jump statement:

  goto L

  where L is label

Common Three-Address Statements (II)

- Conditional jump statement:

  if (x relop y) goto L

  where relop is a relational operator such as =, $\neq$, >, <
- Procedural call statement:

  param $\mathbf{x}_1$, ..., param $\mathbf{x}_n$, call $\mathbf{F}_y$, n

  As an example, foo(x1, x2, x3) is translated to

  param x_1
  param x_2
  param x_3
  call foo, 3
- Procedural call return statement:

  return y

  where y is the return value (if applicable)

Common Three-Address Statements (III)

- Indexed assignment statement:

  x = y[i]

  or

  y[i] = x

  where x is an int and y is an array variable

- Address and pointer operation statement:

  x = & y ; a pointer x is set to the address of y

  y = * x ; y is set to value contained in the location
  ; pointed to by x

  ***y = x** ; location addressed by y gets value x

Static Single Assignment (SSA)

- Developed by R. Cytron, J. Ferrante, *et al.* in 1980s
  - Every variable is assigned exactly once i.e. one **DEF**
  - Convert original variable name to $\text{name}_{version}$
    e.g. $x \rightarrow x_1, x_2$ in different places
  - Use ϕ-function to combine two DEFs of same original variable on a control flow merge

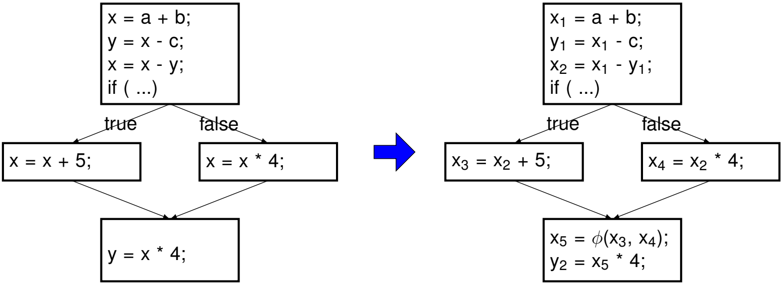

SSA helps compiler optimizations

- Dead Code Elimination (DCE):

$x = a + b;$
$x = c - d;$
$y = x * b;$

$\Rightarrow$

$x_1 = a + b;$
$x_2 = c - d;$
$y_1 = x_2 * b;$

.... x_1 is defined but never used, it is safe to remove

- Partial Redundancy Elimination (PRE):

$x_1 = a_1 + b_1;$
if (...)

true: $a_2 = a_1 * 2;$

false: ...

$a_3 = \phi(a_2, a_1);$
$y_1 = a_3 + b_1;$

$\Rightarrow$

$x_1 = a_1 + b_1;$
if (...)

true: $a_2 = a_1 * 2;$
$t_1 = a_2 + b_1$

false: $t_2 = x_1;$

$t_3 = \phi(t_1, t_2);$
$y_1 = t_3;$

Redundant a + b computation on false branch is removed

Why does SSA help optimizations?

- Data dependencies between instructions are made explicit
  - Variables with same name guaranteed to have same value
- Without SSA, same name does not mean same value
  - Must maintain data dependence graph to express this info
- We will discuss more in **compiler optimization** phase

Laying Out Memory

Layout of Variables in Stack Memory

- Local variables of a function are allocated on stack frame
  - Maintain **offset** from base of frame to allocate next variable
    - address(x) $\leftarrow$ offset
    - offset += sizeof(x.type)

```
void foo() {
    int a;
    int b;
    long long c;
    int d;
}
```

Layout of Variables in Stack Memory

- Local variables of a function are allocated on stack frame
  - Maintain **offset** from base of frame to allocate next variable
    - address(x) ← offset
    - offset += sizeof(x.type)

```
void foo() {
    int a;
    int b;
    long long c;
    int d;
}
```

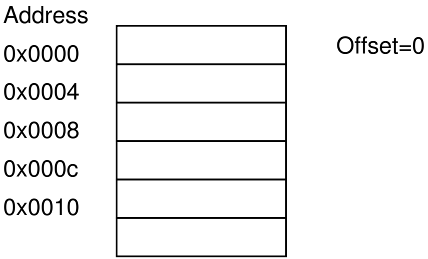

Layout of Variables in Stack Memory

- Local variables of a function are allocated on stack frame
  - Maintain **offset** from base of frame to allocate next variable
    - address(x) ← offset
    - offset += sizeof(x.type)

```
void foo() {
    int a;
    int b;
    long long c;
    int d;
}
```

Address
0x0000
0x0004
0x0008
0x000c
0x0010

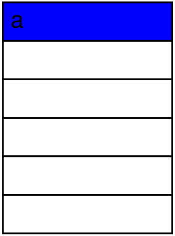

Offset=0
Addr(a)←0

Layout of Variables in Stack Memory

- Local variables of a function are allocated on stack frame
  - Maintain **offset** from base of frame to allocate next variable
    - address(x) ← offset
    - offset += sizeof(x.type)

```
void foo() {
    int a;
    int b;
    long long c;
    int d;
}
```

Address	
0x0000	a
0x0004	
0x0008	
0x000c	
0x0010	

Offset=4
Addr(a)←0

Layout of Variables in Stack Memory

- Local variables of a function are allocated on stack frame
  - Maintain **offset** from base of frame to allocate next variable
    - address(x) $\leftarrow$ offset
    - offset += sizeof(x.type)

```
void foo() {
    int a;
    int b;
    long long c;
    int d;
}
```

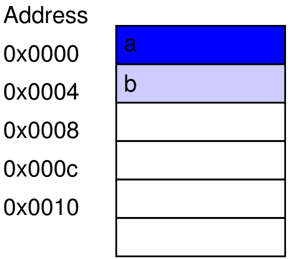

Offset=8
Addr(a)$\leftarrow$0

Addr(b)$\leftarrow$4

Layout of Variables in Stack Memory

- Local variables of a function are allocated on stack frame
  - Maintain **offset** from base of frame to allocate next variable
    - address(x) ← offset
    - offset += sizeof(x.type)

```
void foo() {
    int a;
    int b;
    long long c;
    int d;
}
```

Address	
0x0000	a
0x0004	b
0x0008	c
0x000c	c
0x0010	

Offset=16
Addr(a)←0
Addr(b)←4
Addr(c)←8

Layout of Variables in Stack Memory

- Local variables of a function are allocated on stack frame
  - Maintain **offset** from base of frame to allocate next variable
    - address(x) $\leftarrow$ offset
    - offset += sizeof(x.type)

```
void foo() {
    int a;
    int b;
    long long c;
    int d;
}
```

Address	
0x0000	a
0x0004	b
0x0008	c
0x000c	c
0x0010	d

Offset=20
Addr(a)$\leftarrow$0
Addr(b)$\leftarrow$4
Addr(c)$\leftarrow$8
Addr(d)$\leftarrow$16

What if function has nested scopes?

- Let's take the below example code:

```
void foo() {
    int a;
    int b;
    {
        int i;
    }
    {
        {
            int j;
        }
        int k;
    }
}
```

- What is address(k)? 16?

Nested Scopes Example

```
void foo() {
    int a;
    int b;
      check point #1
    {
        int i;
    }

    {
        {
            int j;
        }
        int k;
    }
}
```

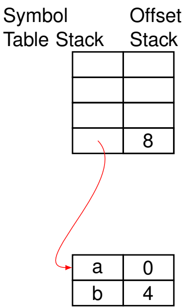

Nested Scopes Example

```
void foo() {
    int a;
    int b;
      check point #1
    {
        int i;
      check point #2
    }

    {
        {
            int j;

        }
        int k;

    }

}
```

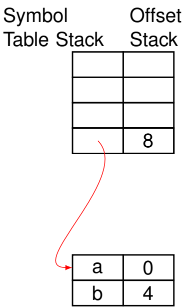

Nested Scopes Example

```
void foo() {
      int a;
      int b;
        check point #1
   {
         int i;
      check point #2
   }

   {
         {
             int j;
         }
         int k;
   }
}
```

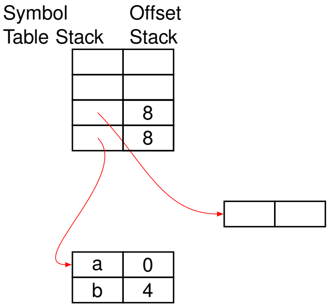

Nested Scopes Example

```
void foo() {
    int a;
    int b;
      check point #1
  {
      int i;
    check point #2
  }

  {
      {
          int j;

      }
      int k;

  }

}
```

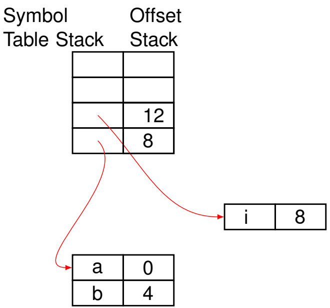

Nested Scopes Example

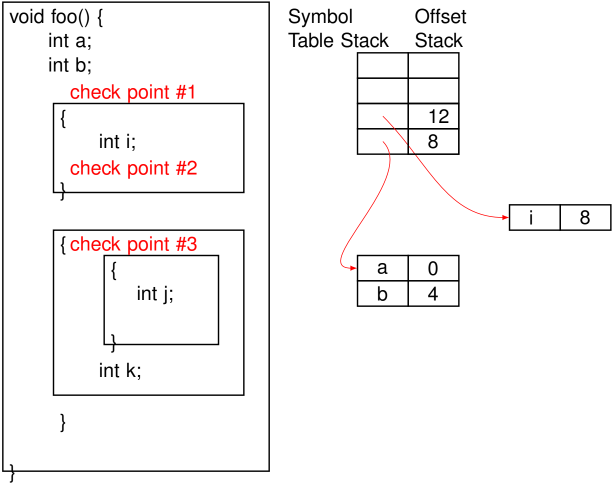

Nested Scopes Example

```
void foo() {
    int a;
    int b;
      check point #1
    {
        int i;
      check point #2
    }

    { check point #3
        {
            int j;

        }
        int k;

    }

}
```

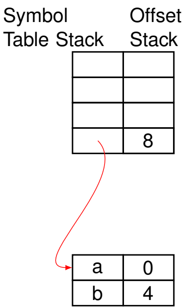

Nested Scopes Example

```
void foo() {
      int a;
      int b;
        check point #1
    {
          int i;
        check point #2
    }

    { check point #3
        {
            int j;

        }
          int k;

    }

}
```

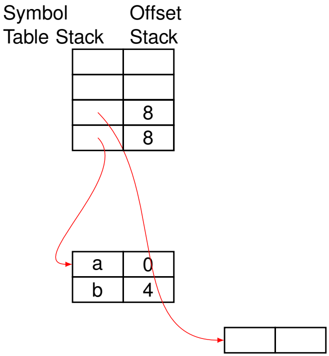

Nested Scopes Example

```
void foo() {
      int a;
      int b;
        check point #1
    {
          int i;
        check point #2
    }

    { check point #3
        {
            int j;
        check point #4
        }
          int k;

    }

}
```

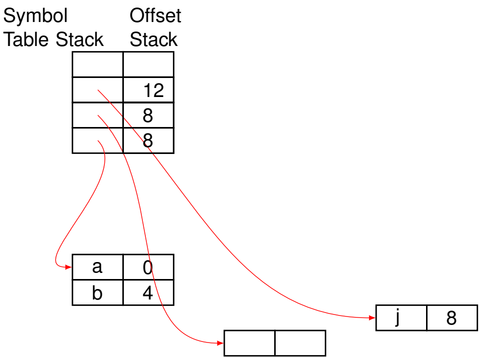

Nested Scopes Example

```
void foo() {
      int a;
      int b;
        check point #1
    {
          int i;
        check point #2
    }

    { check point #3
          {
              int j;
        check point #4
          }
          int k;
        check point #5
    }

}
```

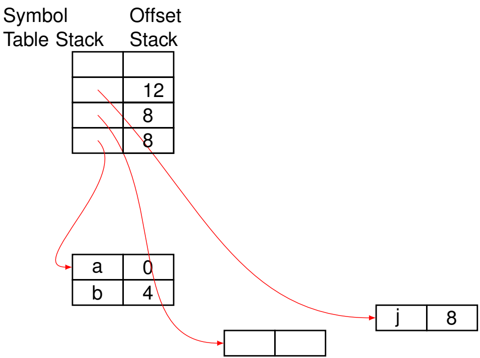

Nested Scopes Example

```
void foo() {
    int a;
    int b;
      check point #1
    {
        int i;
      check point #2
    }

    { check point #3
        {
            int j;
      check point #4
        }
        int k;
      check point #5
    }

}
```

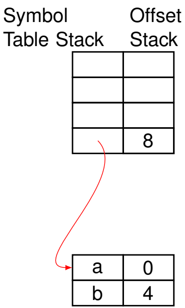

Nested Scopes Example

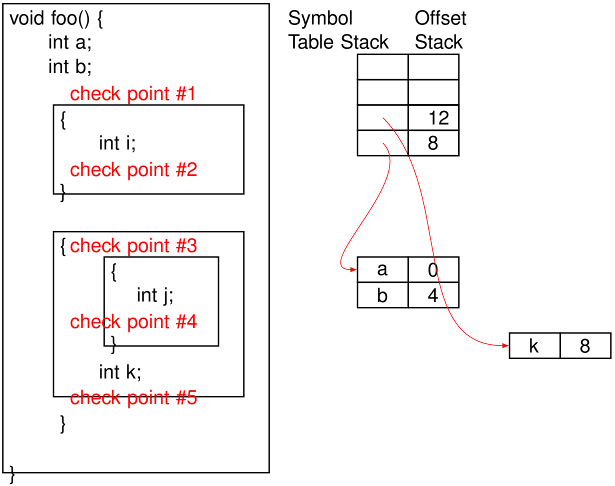

Consideration 1: Allocation Alignment

- Enforce **addr(var) mod sizeof(memory word) == 0**
  - Memory word: unit of memory access in given CPU
  - If not, need to load two words and shift & concatenate

```
void foo() {
    char a;         // addr(a) = 0;
    int b;          // addr(b) = 4; /* instead of 1 */
    int c;          // addr(c) = 8;
    long long d;    // addr(d) = 16; /* instead of 12 */
}
```

- This makes memory layout backend machine dependent
  - Memory layout made explicit only in Machine IR
  - Low-level IR needs to refers to locations in an abstract way

Consideration II: Endianness

- Endianness
  - **Big endian**: **MSB** (most significant byte) in lowest address
  - **Little endian**: **LSB** (least significant byte) in lowest address

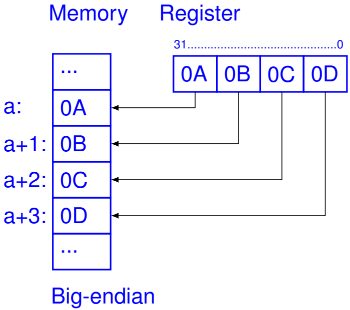

Consideration II: Endianness

- Endianness
  - **Big endian**: **MSB** (most significant byte) in lowest address
  - **Little endian**: **LSB** (least significant byte) in lowest address

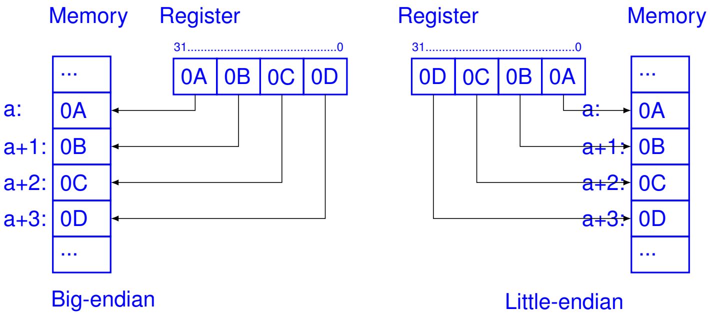

How about other memory besides stack memory?

- Static Memory
  - Where global variables and other static variables reside
  - Layout variables from base address in the same way
- Heap Memory
  - Where dynamically allocated memory using malloc reside
  - Handled by runtime memory management library
  - Compiler not much to do with how this memory is laid out

Generating IR

Generating IR from Language Constructs

- Goal: translate **language constructs** in syntax tree to IR
- Two ways to implement semantic rules for translation
  - By depth-first traversal of syntax tree built by parser
    - This is what you are doing for Project 4
  - By a **syntax directed translation scheme**
    - This is what we will discuss now (based on LR parser)
    - But most concepts apply to both implementations

Generating IR from Language Constructs

- Goal: translate **language constructs** in syntax tree to IR
- Two ways to implement semantic rules for translation
  - By depth-first traversal of syntax tree built by parser
    - This is what you are doing for Project 4
  - By a **syntax directed translation scheme**
    - This is what we will discuss now (based on LR parser)
    - But most concepts apply to both implementations
- What language structures do we need to translate?
  - Declarations
    - Variables, functions (parameters and return types), ...
  - Statements
    - Assignment statements
    - Function call statements
    - Control flow statements (if-then-else, for/while loops)
  - Expressions
    - x + y, x - y, x < y, x > y, x == y, ...

Attributes to Evaluate in Translation

- Variable declaration:
  T V e.g. int a,b,c;
  - Type information **T.type T.width**
  - Variable information **V.type**, **V.offset**

- Statement **S**
  - **S.code**: synthesized attribute that holds IR code of S

- Expression **E**
  - **E.code**: synthesized attribute that holds IR code for E
  - **E.place**: synthesized attribute for temporary variable name to store result of E (for SSA, virtual register name)

Processing Declarations

- Translating declarations in a single scope
  - **enter(name, type, offset)**: insert variable into symbol table

$S \rightarrow M\ D$	
$M \rightarrow \varepsilon$	{ offset=0; } /* reset offset before layout */
$D \rightarrow D\ D$	
$D \rightarrow$ T id;	{ enter(id.name, T.type, offset); offset += T.width; }
$T \rightarrow$ integer	{ T.type=integer; T.width=4; }
$T \rightarrow$ real	{ T.type=real; T.width=8;}
$T \rightarrow$ T1[num]	{ T.type=array(num.val, T1.type); T.width=num.val * T1.width; }
$T \rightarrow$ * T1	{ T.type=ptr(T1.type); T.width=4; }

Processing Declarations in Nested Scopes

- Translating declarations in nested scopes
  - **push(item, stack)**: Pushes item on to stack
  - **pop(stack)**: Pops item at the top of stack
  - **top(stack)**: Returns item at the top of stack

S $\rightarrow$ M D	{ pop(tblptr); pop(offset); }
M $\rightarrow$ ε	{ t=mktable(nil); push(t, tblptr); push(0,offset); }
D $\rightarrow$ D D	
D $\rightarrow$ { N D }	{ pop(tblptr); pop(offset); }
N $\rightarrow$ ε	{ t=mktable(nil); push(t, tblptr); push(top(offset), offset);
D $\rightarrow$ T id;	{ enter(id.name, T.type, top(offset)); top(offset) = top(offset)+ T.width; }

Processing Statements

- Statements rely on symbol table populated by declarations
  - **lookup (id)**: search id in symbol table, return nil if none
  - **emit(code)**: print three address IR for code
  - **newtemp()**: get a new temporary variable (or register)
    - E.g., in SSA form, it returns the next virtual register number
      int newtemp() { return virtual_register++; }

S $\rightarrow$ id = E	{ P=lookup(id); if (P==nil) perror(...); else emit(P '=' E.place);}
E $\rightarrow$ E1 + E2	{ E.place = newtemp(); emit(E.place '=' E1.place '+' E2.place); }
E $\rightarrow$ E1 * E2	{ E.place = newtemp(); emit(E.place '=' E1.place '*' E2.place); }
E $\rightarrow$ - E1	{ E.place = newtemp(); emit(E.place '=' '-' E1.place); }
E $\rightarrow$ (E1)	{ E.place = E1.place; }
E $\rightarrow$ id	{ P=lookup(id); E.place=P; }

Processing Array References

- Recall generalized row/column major addressing

- For example:
  1-dimension: int x[100]; ..... $x[i_1]$
  2-dimension: int x[100][200]; ..... $x[i_1][i_2]$
  3-dimension: int x[100][200][300]; ..... $x[i_1][i_2][i_3]$

- Row major: offset of a k-dimension array item
  1-dimension: $A_1 = a_1$*width $\quad a_1 = i_1$
  2-dimension: $A_2 = a_2$*width $\quad a_2 = a_1{}^*N_2 + i_2$
  3-dimension: $A_3 = a_3$*width $\quad a_3 = a_2{}^*N_3 + i_3$
  ...
  k-dimension: $A_k = a_k$*width $\quad a_k = a_{k-1}{}^*N_k + i_k$

Processing Array References

- Processing an array assignment (e.g. A[i] = B[j];)

S → L = E   { t = newtemp(); emit(t '=' L.place '*' L.width);
    emit(t '=' L.base '+' t); emit ('*'t '=' E.place); }

E → L   { E.place = newtemp(); t= newtemp();
    emit(t '=' L.place '*' L.width); emit (E.place '=' (L.base '+' t)); }

L → id [E]   { L.base = lookup(id).base; L.width = lookup(id).width; L.dim=1;
    L.place = E.place; }

L → L1 [E]   { L.base = L1.base; L.width = L1.width; L.dim = L1.dim + 1;
    L.place = newtemp();
    emit(L.place '=' L1.place '*' L.max[L.dim]);
    emit(L.place '=' L.place '+' E.place); }

Processing Boolean Expressions

- Boolean expression: **a op b**
  - where op can be <, >, >=, &&, ||, ...

1. Without *short circuiting*
   - Short circuiting:
     - In expression A && B, not evaluating B when A is false
     - In expression A || B, not evaluating B when A is true
   - Without short circuiting, entire expression is evaluated:

$S \rightarrow$ id = E	$\equiv$	lookup(id) = E.place
$E \rightarrow$ (a < b) or (c < d and e < f)	$\equiv$	t1 = a < b
		t2 = c < d
		t3 = e < f
		t4 = t2 && t3
		E.place = t1 || t4

Processing Boolean Expressions

2. With short circuiting (e.g. C/C++/Java)
   - ➢ Processing simple boolean expressions:

     S→ if E then S1

     E→ a < b ≡ E.true = code for S1;
     E.false = S.next;
     emit(if E goto E.true);
     emit(goto E.false);

     S.next: address of code after S

     E.true: address of code to execute on 'true'

     E.false: address of code to execute on 'false'
   - ➢ Processing compound boolean expressions:
     - Chain together multiple of above by updating E.true/E.false
     - E→ E1 && E2: E1.true = code for E2, E1.false = S.next
     - E→ E1 || E2: E1.false = code for E2, E1.true = code for S1

Processing Boolean Expressions

2. With short circuiting (cont'd)
   - A short circuited compound boolean expression
     S→ if E then S1
     E→ (a < b) or (c < d and e < f) ≡

```
     E.true = code for S1;
     E.false = S.next;
     emit(
     if (a<b) goto E.true
     goto L1
L1:  if (c<d) goto L2
     goto E.false
L2:  if (e<f) goto E.true
     goto E.false
     );
```

   - **Problem: E.true, E.false, S.next are non-L-attributes**
     - Depend on code that has not been generated yet
       E.true: Only available when S1 is generated
       E.false: Only available when code after S is generated
     - Emitting any **forward jump** poses this problem

Syntax Directed Translation

- Non-L-attributs preclude syntax directed translation
  - Even with parse tree, preclude simple left-to-right traversal
- Solutions: two methods
  - Two pass approach — process the code twice
    - Generate labels in the first pass
    - Replace labels with addresses in the second pass
  - One pass approach
    - Generate holes when address is needed but unknown
    - Backpatch holes when address is known later on
- We will discuss the more efficient one pass approach
  - It is also the method you will use in project 4.

One-Pass Based Syntax Directed Translation

- Non-L-attributes during code generation are unavoidable
  - Due to forward jumps to code on the right hand side
  - Example: E.true, E.false, S.next for boolean expressions
  - Is there a one-pass solution to the problem?

Idea:

1. Leave holes for non-L-attribute values we don't know
2. Fill the holes in when we know the values later on
   - *holelist*: synthesized attribute of 'holes' for one value
   - Holes are filled in by traversing list when value is known
   - All holes will be patched by the end of code generation (Since all forward jumps would be resolved by then)

One-Pass Based Syntax Directed Translation

- Attributes for two pass based approach
  - Expression **E**
    —- Synthesized attributes: **E.code** **E.holes_truelist**, and **E.holes_falselist**
  - Statement **S**
    —- Synthesized attributes: **S.code** and **S.holes_nextlist**

- 3 functions for implementing backpatching
  - **makelist(i)**: create a holelist with statement index i
  - **merge(p1, p2)**: concatenate list p1 and list p2
  - **backpatch(p, i)**: insert index i in every hole in holelist p

Backpatching for if-then

- Given rule **S $\rightarrow$ if E then S1**, below is done in one-pass:
  (1). Gen **E.code**, making **E.holes_truelist**, **E.holes_falselist**
  (2). Gen **S1.code**, filling in **E.holes_truelist** and merging **S1.holes_nextlist** with **E.holes_falselist**
  (3). Pass on merged list to **S.holes_nextlist**

S $\rightarrow$ if E then M S1
{
backpatch(E.holes_truelist, M.index);
S.holes_nextlist = merge(S1.holes_nextlist, E.holes_falselist);
}

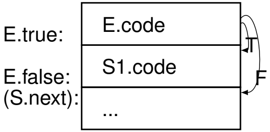

M $\rightarrow \varepsilon$
{ M.index = curIndex; }

Backpatching for if-then-else

- Given rule **S** $\rightarrow$ **if E then S1 else S2**:
  (1). Gen **E.code**, making **E.holes_truelist**, **E.holes_falselist**
  (2). Gen **S1.code**, filling in **E.holes_truelist**
  (3). Emit goto to S.next before S2, to skip over S2
  (4). Gen **S2.code**, filling in **E.holes_falselist**
  (5). Merge relevant holes into **S.holes_nextlist**

S $\rightarrow$ if E then M1 S1 N else M2 S2

```
{
backpatch(E.holes_truelist, M1.index);
backpatch(E.holes_falselist, M2.index);
templist = merge(S1.holes_nextlist, N.holes_nextlist);
S.holes_nextlist = merge(templist, S2.holes_nextlist);
}
```

N $\rightarrow$ ε

```
{
N.holes_nextlist = makelist(curIndex);
emit('goto ___');
}
```

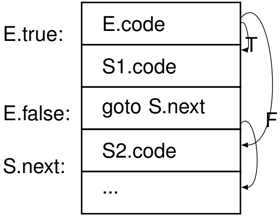

Backpatching for S.holes_nextlist

- When does the holes in S.holes_nextlist get patched?

- When the instruction after S is generated of course!

- Given rule **S** $\rightarrow$ **S1 S2**:
  (1). Gen **S1.code** making **S1.holes_nextlist**
  (2). Gen **S2.code** filling in **S1.holes_nextlist** and making **S2.holes_nextlist**
  (3). Pass on **S2.holes_nextlist** to **S.holes_nextlist**

Backpatching for Boolean Expressions

E $\rightarrow$ E1 or M E2	{ backpatch(E1.holes_falselist, M.index); E.holes_truelist = merge(E1.holes_truelist, E2.holes_truelist); E.holes_falselist = E2.holes_falselist; }
E $\rightarrow$ E1 and M E2	{ backpatch(E1.holes_truelist, M.index); E.holes_falselist = merge(E1.holes_falselist, E2.holes_falselist); E.holes_truelist = E2.holes_truelist; }
M $\rightarrow$ ε	{ M.index = curIndex; }

Backpatching for Boolean Expressions

E → not E1	{ E.holes_truelist = E1.holes_falselist; E.holes_falselist = E1.holes_truelist; }
E → (E1)	{ E.holes_truelist = E1.holes_truelist; E.holes_falselist = E1.holes_falselist; }
E → id1 relop id2	{ E.holes_truelist = makelist(curIndex); E.holes_falselist = makelist(curIndex+1); emit('if' id1.place 'relop' id2.place 'goto ___'); emit('goto ___'); }
E → true	{ E.holes_truelist = makelist(curIndex); emit('goto ___'); }
E → false	{ E.holes_falselist = makelist(curIndex); emit('goto ___'); }

Backpatching Example

- E $\rightarrow$ (a<b) or M1 (c<d and M2 e<f)

- When reducing (a<b) to E1, we have
  100: if(a<b) goto ___ E1.hole_truelist=(100)
  101: goto ___ E1.hole_falselist=(101)
- When reducing ε to M1, we have M1.index = 102
- When reducing (c<d) to E2, we have
  102: if(c<d) goto ___ E2.hole_truelist=(102)
  103: goto ___ E2.hole_falselist=(103)
- When reducing ε to M2, we have M2.index = 104
- When reducing (e<f) to E3, we have
  104: if(e<f) goto ___ E3.hole_truelist=(104)
  105: goto ___ E3.hole_falselist=(105)

Backpatching Example (cont.)

- When reducing (E2 and M2 E3) to E4, we backpatch((102), 104);

100: if(a<b) goto ___	E4.hole_truelist=(104)
101: goto ___	E4.hole_falselist=(103,105)
102: if(c<d) goto **104**	
103: goto ___	
104: if(e<f) goto ___	
105: goto ___	

- When reducing (E1 or M1 E4) to E5, we backpatch((101), 102);

100: if(a<b) goto ___	E5.hole_truelist=(100, 104)
101: goto **102**	E5.hole_falselist=(103,105)
102: if(c<d) goto 104	
103: goto ___	
104: if(e<f) goto ___	
105: goto ___	

Backpatching Example (cont.)

- When reducing (E2 and M2 E3) to E4, we backpatch((102), 104);

```
100: if(a<b) goto ___            E4.hole_truelist=(104)
101: goto ___                    E4.hole_falselist=(103,105)
102: if(c<d) goto 104
103: goto ___
104: if(e<f) goto ___
105: goto ___
```

- When reducing (E1 or M1 E4) to E5, we backpatch((101), 102);

```
100: if(a<b) goto ___            E5.hole_truelist=(100, 104)
101: goto 102                    E5.hole_falselist=(103,105)
102: if(c<d) goto 104
103: goto ___
104: if(e<f) goto ___
105: goto ___
```

- Are we done?

Backpatching Example (cont.)

- When reducing (E2 and M2 E3) to E4, we backpatch((102), 104);

```
100: if(a<b) goto ___          E4.hole_truelist=(104)
101: goto ___                  E4.hole_falselist=(103,105)
102: if(c<d) goto 104
103: goto ___
104: if(e<f) goto ___
105: goto ___
```

- When reducing (E1 or M1 E4) to E5, we backpatch((101), 102);

```
100: if(a<b) goto ___          E5.hole_truelist=(100, 104)
101: goto 102                  E5.hole_falselist=(103,105)
102: if(c<d) goto 104
103: goto ___
104: if(e<f) goto ___
105: goto ___
```

- Are we done?
  - Yes for this expression

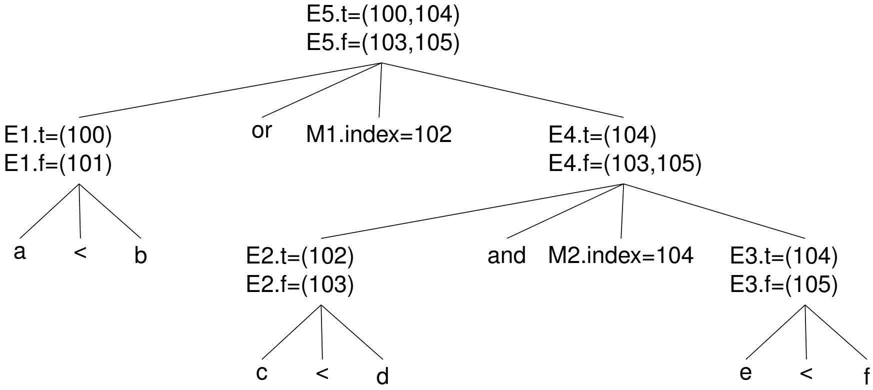
E5.t=(100,104)
E5.f=(103,105)
E1.t=(100)
E1.f=(101)
or
M1.index=102
E4.t=(104)
E4.f=(103,105)
a
<
b
E2.t=(102)
E2.f=(103)
and
M2.index=104
E3.t=(104)
E3.f=(105)
c
<
d
e
<
f