

# Lexical Analysis

# What is Lexical Analysis

- ❑ Comes from the word lexicon, or dictionary
- ❑ **Lexical Analysis:** Partitioning a string into words
  - These words are also called **tokens**
- ❑ We will use this code as a running example:

```
if (i==j)
    z = 0;
else
    z = 1;
```
- ❑ Code is provided as input string to lexical analysis
  - `"if(i == j)\n\tz = 0; \nelse\n\tz = 1; \n"`
- ❑ Goal is turning the string into tokens, or tokenization

# What is a Token ?

- ❑ Smallest unit that has meaning in a string
  - In English, tokens are English words:  
nouns, verbs, adjectives, ...
  - In a programming language:  
identifiers, integers, keywords, whitespace, ...
  
- ❑ A token is a tuple (`type`, `lexeme`)
  - `type`: the token type that the token belongs to
    - Identifier: string of letters and digits, starting with a letter
    - Integer: string of digits
    - Keyword: “else”, “if”, “while”, ...
    - Whitespace: string of blanks, newlines, and tabs
  - `lexeme`: actual string value of this token

# Lexical Analysis is the act of Tokenization

- ❑ Output of lexical analysis is a stream of tokens
- ❑ Tokens are the input to Syntax Analysis (a.k.a. Parsing)
  - Parsers rely on token type to figure out role of each token  
E.g. a keyword is treated differently from an identifier

# Lexical Analysis Tokenization Example

❑ Given `"if(i == j)\n\tz = 0; \tnelse\n\tz = 1; \n"`

❑ What would be an output of lexical analysis?

Recall a token is a tuple `(type, lexeme)`

❑ Output:

`(keyword, "if")``(left-parenthesis, "(")``(identifier, "i")``(equals-op, "=")``(identifier, "j")``(right-parenthesis, ")")``(whitespace, "\n\t")``(identifier, "z")``(assign-op, "=")``(integer, "0")``(semicolon, ";")``(whitespace, "\n")``(keyword, "else")``(whitespace, "\n\t")``(identifier, "z")``(assign-op, "=")``(integer, "1")``(semicolon, ";")``(whitespace, "\n")`

❑ The lexer usually discards “non-interesting” tokens that don’t contribute to parsing, e.g., whitespace, comments

# Some language features makes lexing difficult

- ❑ FORTRAN compilation rule: **whitespace is insignificant**
  - Reason: inaccuracy of card punching by operators
- ❑ Consider
  - DO 5I=1,25
  - DO 5I=1.25
- ❑ This is the interpretation of the two statements:
  - Former: Iterate from I=1 to I=25 with step size 5
  - Latter: Assign 1.25 to variable DO5I
- ❑ Reading left-to-right, cannot tell if DO5I is a variable or DO statement; Have to continue until “,” or “.” is reached.
  - “lookahead” may be required to decide on tokens
  - Feedback necessary from parser to lexical analysis

# C++ language has difficult features too

- ❑ C++ has the Right Angle Brackets >> issue:  
<https://www.open-std.org/jtc1/sc22/wg21/docs/papers/2005/n1757.html>
- ❑ `typedef std::vector<std::vector<int> > Table; // OK`  
`typedef std::vector<std::vector<int>> Table; // Error`
- ❑ Why? >> is read as one token which can either be:
  - A stream operator (e.g. `cin >> var`)
  - Or a shift operator (e.g. `1 >> 2`)
- ❑ Space is needed in between to create two > tokens
- ❑ Fixed in C+11 standard so this is no longer an error
  - That makes tokenization decision on >> context dependent
  - Again forcing lexical analysis to get feedback from parser

# Lexical Analysis Implementation

## Step 1:

- Define a set of token types
  - Refer to language specifications
  - Types you choose depends on design of parser
  - Recall “ $if(i == j)\backslash n\backslash tz = 0; \backslash nelse\backslash n\backslash tz = 1; \backslash n$ ”
  - Should “==” be one token? or two tokens? Depends.
  - Should “if”, “then”, “else” be separate types or just one keyword type? Depends.

## Step 2:

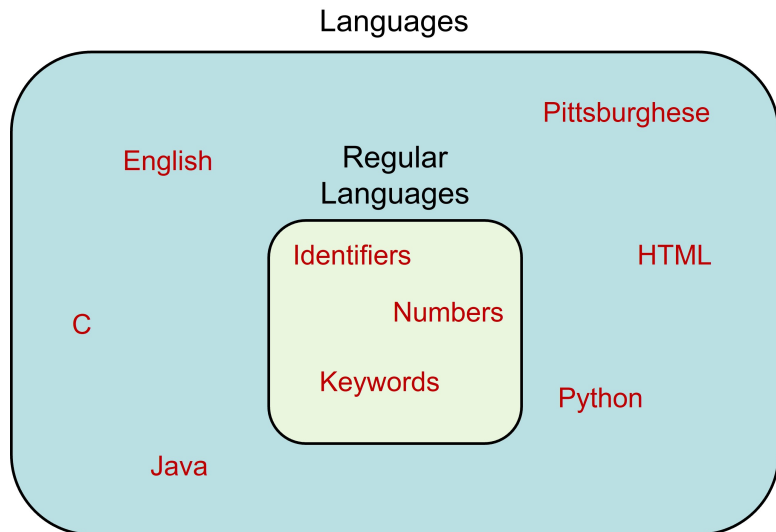
- For each token type, describe which string belongs to it



# Describing strings belonging to a token type

- ❑ A token type is something that looks like this:  
Identifier: string of letters and digits, starting with a letter
- ❑ Is there a more formal (mathematical) way to express this?
  - Yes! By using the formalism of **Languages**.
- ❑ Definition of **Language**:  
Let  $\Sigma$  be a set of characters, a **language** over  $\Sigma$  is a set of strings of the characters drawn from  $\Sigma$ 
  - So by definition, any token type is a Language
  - And so are English, Java, Python, and HTML
- ❑ Some Languages can be very difficult to express formally
  - Imagine having to formally describe the English language!

# Token types belong to a subset of Languages (called Regular Languages)




# Regular Expressions express Regular Languages

## Definition of **Regular Expression**

The **regular expressions (REs)** over  $\Sigma$  are the total set of expressions that can be constructed using the following components:

- $\epsilon$
- 'c' where  $c \in \Sigma$
- $A + B$  where  $A, B$  are **RE** over  $\Sigma$
- $AB$  where  $A, B$  are **RE** over  $\Sigma$
- $A^*$  where  $A$  is a **RE** over  $\Sigma$

 **Regular Languages** are defined as languages that can be expressed using **Regular Expressions**.

# Atomic Regular Expressions

- Single character denotes a set of one string  
 $'c' = \{ "c" \}$
- Epsilon* or  $\epsilon$  character denotes a zero length string  $\epsilon = \{ "" \}$
- Empty set is  $\{ \} = \phi$ , not the same as  $\epsilon$   
 $\text{size}(\phi) = 0$   
 $\text{size}(\epsilon) = 1$   
 $\text{length}(\epsilon) = 0$

# Compound Regular Expressions

- Union: if A and B are REs, then

$$A + B = \{ s \mid s \in A \text{ or } s \in B \}$$

- Concatenation of sets/strings

$$AB = \{ ab \mid a \in A \text{ and } b \in B \}$$

- Iteration (Kleene closure)

$$A^* = \bigcup_{i \geq 0} A^i \quad \text{where } A^i = A \dots A \text{ (} i \text{ times)}$$

in particular

$$A^* = \{ \epsilon \} + A + AA + AAA + \dots$$

$$A^+ = A + AA + AAA + \dots = A A^*$$


# The $L(\textit{Expression})$ Notation

□  $L(\textit{Expression})$ : means the language defined by *Expression*

□ Some example languages defined using the  $L()$  notation:

- $L(\epsilon) = \{ \epsilon \}$
- $L('c') = \{ c \}$
- $L(A+B) = L(A) \cup L(B)$
- $L(AB) = \{ ab \mid a \in L(A) \text{ and } b \in L(B) \}$
- $L(A^*) = \bigcup_{i \geq 0} L(A^i)$

# Examples

 Keywords: “else” or “if” or “while” or ...

# Examples

Keywords: “else” or “if” or “while” or ...

➤ ‘else’ + ‘if’ + ‘while’ + ...

➤ ‘else’ abbreviates

‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’



# Examples



Keywords: “else” or “if” or “while” or ...

➤ ‘else’ + ‘if’ + ‘while’ + ...

➤ ‘else’ abbreviates

‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’

➤ keywords = { ‘else’, ‘if’, ‘then’, ‘while’, ... }

# Examples

## Keywords: “else” or “if” or “while” or ...

- ‘else’ + ‘if’ + ‘while’ + ...
- ‘else’ abbreviates  
‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’
- keywords = { ‘else’, ‘if’, ‘then’, ‘while’, ... }

## Integer

- digit = ‘0’ + ‘1’ + ‘2’ + ‘3’ + ‘4’ + ‘5’ + ‘6’ + ‘7’ + ‘8’ + ‘9’
- integer = digit digit\* = digit<sup>+</sup>

# Examples

## Keywords: “else” or “if” or “while” or ...

- ‘else’ + ‘if’ + ‘while’ + ...
- ‘else’ abbreviates  
‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’
- keywords = { ‘else’, ‘if’, ‘then’, ‘while’, ... }

## Integer

- digit = ‘0’ + ‘1’ + ‘2’ + ‘3’ + ‘4’ + ‘5’ + ‘6’ + ‘7’ + ‘8’ + ‘9’
- integer = digit digit\* = digit<sup>+</sup>
  - **Q:** Is ‘000’ an integer?

# Examples

## Keywords: “else” or “if” or “while” or ...

- ‘else’ + ‘if’ + ‘while’ + ...
- ‘else’ abbreviates  
‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’
- keywords = { ‘else’, ‘if’, ‘then’, ‘while’, ... }

## Integer

- digit = ‘0’ + ‘1’ + ‘2’ + ‘3’ + ‘4’ + ‘5’ + ‘6’ + ‘7’ + ‘8’ + ‘9’
- integer = digit digit\* = digit<sup>+</sup>
  - **Q:** Is ‘000’ an integer?
  - **Q:** Define an RE that excludes the above sequence.

# More Examples

- ❑ Identifier: strings of letters or digits, starting with a letter
  - letter = 'A' + ... + 'Z' + 'a' + ... + 'z'
  - Identifier = letter (letter + digit)\*
  
- ❑ Whitespace: a non-empty string of blanks, newlines, tabs
  - whitespace = ( ' ' + '\n' + '\t' ) +

# More Examples

❏ Identifier: strings of letters or digits, starting with a letter

- letter = 'A' + ... + 'Z' + 'a' + ... + 'z'
- Identifier = letter (letter + digit)\*
  - **Q:** is letter (letter\* + digit\*) the same?

❏ Whitespace: a non-empty string of blanks, newlines, tabs

- whitespace = ( ' ' + '\n' + '\t' ) +

# More Examples

☐ Phones number: consider (412) 624-0000

- $\Sigma = \text{digit} \cup \{ -, (, ) \}$
- $\text{area} = \text{digit}^3$
- $\text{exchange} = \text{digit}^3$
- $\text{phone} = \text{digit}^4$
- $\text{phoneNumber} = '(' \text{ area } ')' \text{ exchange } '-' \text{ phone}$

☐ Email address: student @ pitt.edu

- $\Sigma = \text{letter} \cup \{ ., @ \}$
- $\text{name} = \text{letter}^+$
- $\text{emailAddress} = \text{name } '@' \text{ name } '.' \text{ name}$

# More Regular Expression Notations

□ Some “syntactic sugar” for regular expressions:

Union:  $A + B \equiv A \mid B$

Option:  $A + \varepsilon \equiv A ?$

Range:  $'a' + 'b' + \dots + 'z' \equiv [a-z]$

Excluded range:  
complement of  $[a-z] \equiv [^a-z]$



# Implementation of Lexical Analysis

- We learnt how to precisely define token types
  - Regular expression (RE)

# Implementation of Lexical Analysis

- We learnt how to precisely define token types
  - Regular expression (RE)
- How do we get from the RE to the actual tokenizer?


# Implementation of Lexical Analysis

- We learnt how to precisely define token types
  - Regular expression (RE)
- How do we get from the RE to the actual tokenizer?
  - **Solution 1:** Write the code manually from REs

# Implementation of Lexical Analysis

- ❑ We learnt how to precisely define token types
  - Regular expression (RE)
  
- ❑ How do we get from the RE to the actual tokenizer?
  - **Solution 1:** Write the code manually from REs
  
  - **Solution 2:** Convert RE to code using a tool
    - Tool generates the source code from the given REs
    - Lex (for C), Flex (for C++), JLex (for Java)

# Lex: a Tool for Lexical Analysis



- ❑ Lex enables token definition using declarative coding
  - Declare tokens using regular expressions
  - Write action performed for each RE using imperative coding
- ❑ We will first describe structure of specification file
- ❑ The internals of the tool will be discussed later

# Example Lex Specification File

```
/* 1. Regular expression definitions section */
%{
/* Code block inserted for includes and declarations */
#include <stdlib.h>
}%
string    [a-z]+
space     [ ]+
%%

/* 2. Rules section: action for each regular expression */
{string}  { printf("lexeme: %s, len=%d\n", yytext, yyleng); }
{space}   { /* No action */ }
%%

/* 3. User code section */
int main() {
    while( yylex() != 0 ) {}
    return 0;
}
```

# Example Lex Specification File: Explanation

## Overview of operation:

- 1 Parser calls `yylex()` when ready to process the next token
- 2 `yylex()` tokenizes longest string that matches an RE
- 3 `yylex()` stores the token lexeme in `yytext`
- 4 `yylex()` stores length of lexeme in `yyleng`
- 5 `yylex()` executes the action `{ ... }` in the rule for RE
- 6 If action returns a value, `yylex()` returns that value.  
If action doesn't return, `yylex()` scans next token.  
`yylex()` returns 0 if no more tokens (EOF).

## To test the lexer without a parser, we need a lexer driver

- The `main()` function serves as the lexer driver
- On piping "hello world!" to input of lexer (a.out):

```
$ echo "hello world!" | ./a.out
lexeme:  hello, len=5
lexeme:  world, len=5
!
```

# How is the Specification File Converted to a Lexer

□ The problem we face is

Given a string **s** and a regular expression **RE**, is

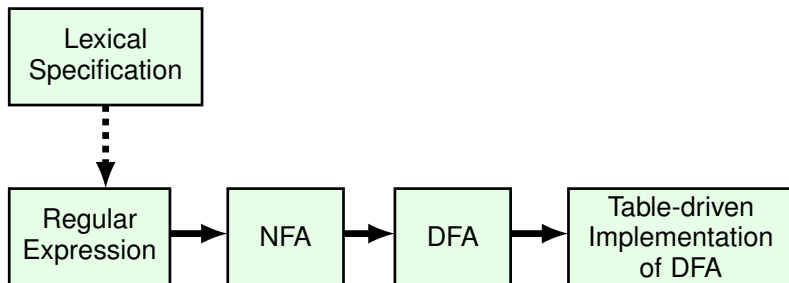
$$\mathbf{s} \in \mathbf{L(RE)} ?$$



# Implementing Lexical Analysis with Finite Automata

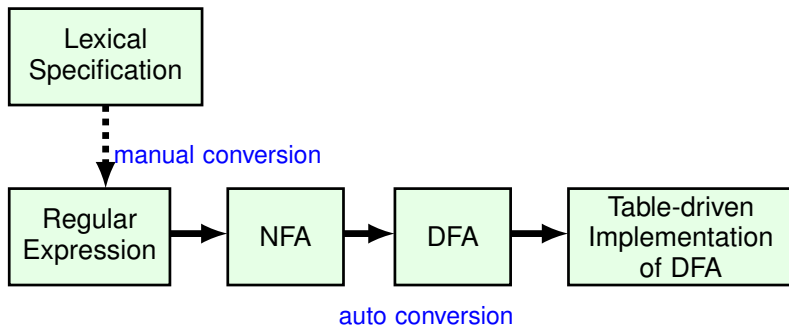
# An Overview of RE to FA

## Our implementation sketch



# An Overview of RE to FA

## Our implementation sketch



# Implementation Outline

## RE $\Rightarrow$ NFA $\Rightarrow$ DFA $\Rightarrow$ Table-driven Implementation

- Tokens are specified using Regular Expressions
- Tokens are accepted using a table-driven DFA
  - Deterministic Finite Automata (DFAs)
  - Non-deterministic Finite Automata (NFAs)
- Table implementations

## I will soon show $RL \equiv L(RE) \equiv L(NFA) \equiv L(DFA)$

- Will show an automated way to take RE all the way to DFA

# Finite Automata


□ A finite automaton consists of 5 components  
 $(\Sigma, S, n, F, \delta)$

- (1). An input alphabet  $\Sigma$
- (2). A set of states  $S$
- (3). A start state  $n \in S$
- (4). A set of accepting states  $F \subseteq S$
- (5). A set of transitions  $\delta: S_a \xrightarrow{\text{input}} S_b$

# State Graph is a graphical representation of FA

- ☐ Sometimes we use **state graph** to represent a FA
- ☐ A **state graph** includes

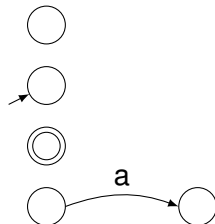
- A set of states
- A start state
- A set of accepting states
- A set of transitions



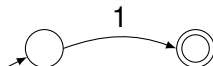
# State Graph is a graphical representation of FA

- Sometimes we use **state graph** to represent a FA
- A **state graph** includes

- A set of states
- A start state
- A set of accepting states
- A set of transitions

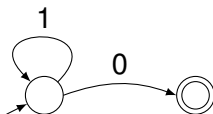


- Example: a finite state automaton that accepts only "1"



# More State Graph Examples


- A finite automaton accepting any number of **1**s followed by a single **0**. Here we have Alphabet =  $\{0,1\}$





# More State Graph Examples

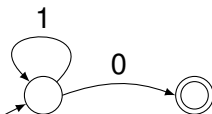
- A finite automaton accepting any number of **1**s followed by a single **0**. Here we have Alphabet =  $\{0,1\}$




- Example: What language does the following state graph recognize? Here we have Alphabet =  $\{0,1\}$

# More State Graph Examples

- A finite automaton accepting any number of **1**s followed by a single **0**. Here we have Alphabet =  $\{0,1\}$



- Example: What language does the following state graph recognize? Here we have Alphabet =  $\{0,1\}$



# How tokens are accepted by Finite Automata


Transition  $\delta: S_a \xrightarrow{c} S_b$  read as:

When in state  $S_a$ , go to state  $S_b$  when scanned char is “c”

- 1 Begin at start state  $n \in S$ .
- 2 Transition until end of input or no transition possible
- 3 Read current state  $x$ 
  - If  $x \in$  accepting set  $F$ , then  $\Rightarrow$  **accept**
  - Otherwise,  $\Rightarrow$  **reject**

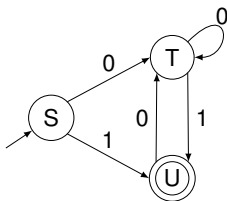
# Table Implementation of a DFA

Given the state graph of a DFA,



# Table Implementation of a DFA

Given the state graph of a DFA,




→ input characters

state ↓

	0	1
S		
T		
U		

# Table Implementation of a DFA

Given the state graph of a DFA,



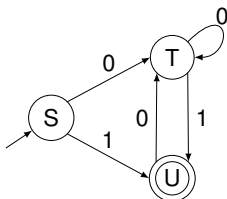
→ input characters

state ↓

	0	1
S	T	U
T	T	U
U	T	X

# Table Implementation of a DFA

Given the state graph of a DFA,



→ input characters

state ↓

	0	1
S	T	U
T	T	U
U	T	X

## Table-driven Code:

```

DFA() {
    state = "S";
    while (!done) {
        ch = fetch_input();
        state = Table[state][ch];
        if (state == "x")
            perror("error");
    }
    if (state ∈ F)
        printf("accept");
    else
        printf("reject");
}
  
```

# Table-driven Code is identical across DFAs

- ❑ Each RE has a different DFA, meaning different table
- ❑ But table-driven code for tokenization remains the same!



# Table-driven Code is identical across DFAs

- Each RE has a different DFA, meaning different table

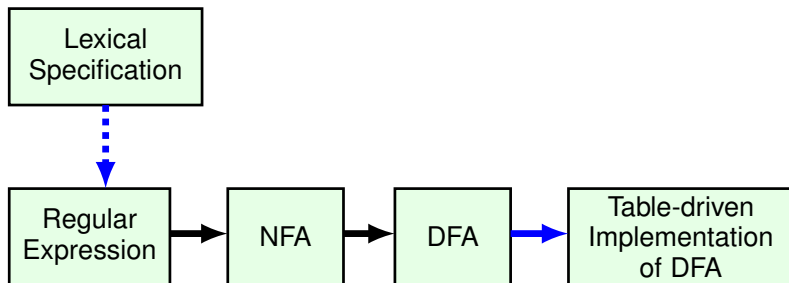
- But table-driven code for tokenization remains the same!

- Revisit our implementation outline

RE  $\Rightarrow$  NFA  $\Rightarrow$  **DFA**  $\Rightarrow$  **Table-driven Implementation**

# From RE to FA


## Our implementation sketch



# Epsilon Moves

## ■ Another kind of transition: $\epsilon$ -moves

- Moves from state A to state B without reading input




## ■ $\epsilon$ -moves makes transitions nondeterministic

- Instead of next character determining the next state, now machine has a choice of staying in A or moving to B

# Deterministic and Nondeterministic Automata

- ❑ Deterministic Finite Automata (DFA)
  - One transition per input per state
  - No  $\epsilon$ -moves
- ❑ Non-deterministic Finite Automata (NFA)
  - Can have multiple transitions for one input in a given state
  - Can have  $\epsilon$ -moves
- ❑ DFAs are easier to implement
  - Real machines are deterministic by nature just like DFAs
  - Given a nondeterministic choice, hard to know what to do

# Examples




# Converting RE to NFA


## McNaughton-Yamada-Thompson Algorithm

### Step 1: processing atomic REs

➤  $\epsilon$  expression



➤ single character RE a



# Converting RE to NFA (cont.)

## Step 2: processing compound REs

➤  $r = s \mid t$


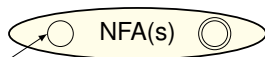
➤  $r = s t$

➤  $r = s^*$

# Converting RE to NFA (cont.)

## Step 2: processing compound REs

➤  $r = s \mid t$



➤  $r = s t$

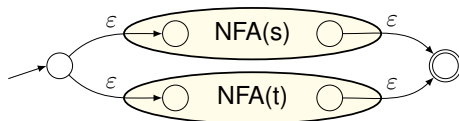
➤  $r = s^*$



# Converting RE to NFA (cont.)

## Step 2: processing compound REs

➤  $r = s \mid t$



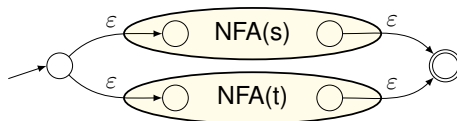
➤  $r = s t$

➤  $r = s^*$

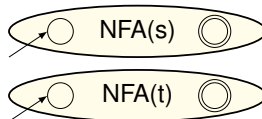
# Converting RE to NFA (cont.)

## Step 2: processing compound REs

➤  $r = s \mid t$



➤  $r = s t$

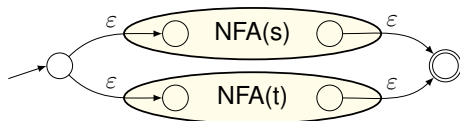


➤  $r = s^*$

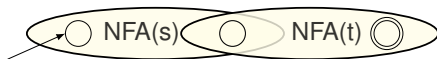
# Converting RE to NFA (cont.)

## Step 2: processing compound REs

➤  $r = s \mid t$



➤  $r = s t$




➤  $r = s^*$


# Converting RE to NFA (cont.)

## Step 2: processing compound REs


➤  $r = s \mid t$



➤  $r = s t$



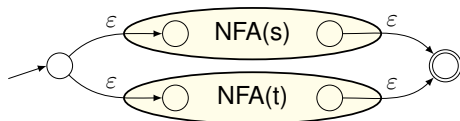
➤  $r = s^*$



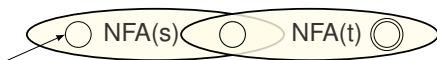
# Converting RE to NFA (cont.)

## Step 2: processing compound REs

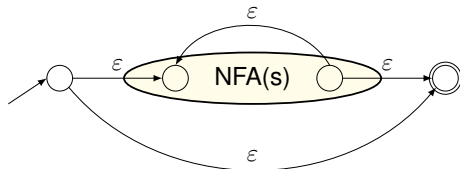
➤  $r = s \mid t$




➤  $r = s t$



➤  $r = s^*$




# In-class Practice

 Convert “**(a|b)\*a b b**” to NFA

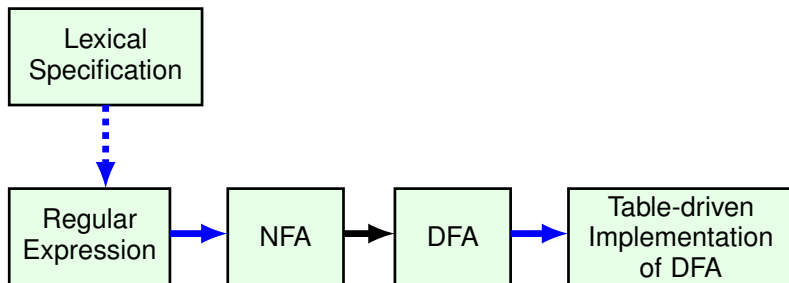
# In-class Practice

Convert “**(a|b)\* a b b**” to NFA



# From RE to FA

## Our implementation sketch





# Execution of Finite Automata

- ❑ A DFA can take only one path through the state graph
  - Completely determined by input
- ❑ A NFA has a choice of:
  - Whether to make  $\epsilon$ -moves
  - Which of multiple transitions for a single input to take

# Execution of Finite Automata



- ❑ A DFA can take only one path through the state graph
  - Completely determined by input
- ❑ A NFA has a choice of:
  - Whether to make  $\epsilon$ -moves
  - Which of multiple transitions for a single input to take
- ❑ When to accept a token for DFAs and NFAs
  - DFA: When input string leads to a final state
  - NFA: When input string **\*\*can\*\*** lead to a final state

# Execution of Finite Automata

- ❑ A DFA can take only one path through the state graph
  - Completely determined by input
- ❑ A NFA has a choice of:
  - Whether to make  $\epsilon$ -moves
  - Which of multiple transitions for a single input to take
- ❑ When to accept a token for DFAs and NFAs
  - DFA: When input string leads to a final state
  - NFA: When input string **\*\*can\*\*** lead to a final state
- ❑ **Question:** Which one is more expressive?

$$L(NFA) \equiv L(DFA) \equiv RL$$

Both accept “**(a|b)\* a b b**”



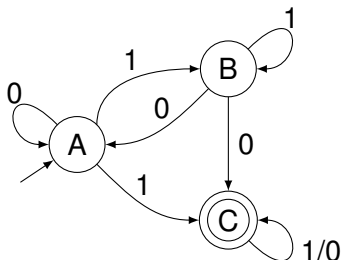
# How to Convert NFA to DFA

- ❑ Basic idea: Given NFA, simulate its execution using DFA
  - At step  $n$ , the NFA may be in any of multiple possible states
- ❑ The new DFA is constructed as follows,
  - A state of DFA  $\equiv$  a non-empty subset of states of the NFA
  - Start state  $\equiv$  the set of NFA states reachable through  $\epsilon$ -moves from NFA start state
  - A transition  $S_a \xrightarrow{c} S_b$  is added **iff**

$S_b$  is the set of NFA states reachable from any state in  $S_a$  after seeing the input  $c$ , considering  $\epsilon$ -moves as well

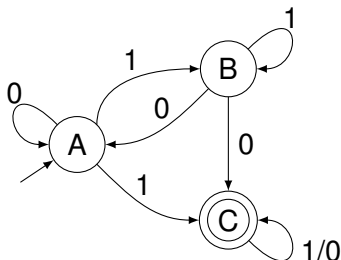
# Example NFA to DFA

What is the Equivalent DFA ?



# Example NFA to DFA

What is the Equivalent DFA ?

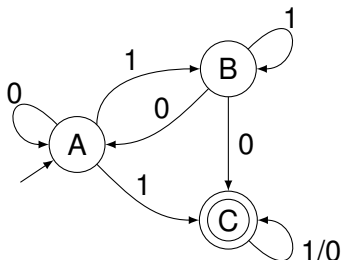


state ↓ → input characters

	0	1
A		
B		
C		

# Example NFA to DFA

What is the Equivalent DFA ?



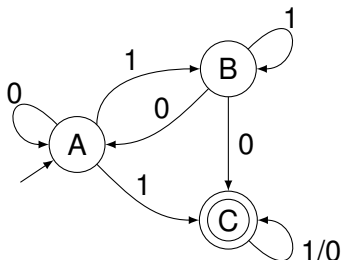
state ↓ → input characters

	0	1
A	A	BC
B	AC	B
C	C	C



# Example NFA to DFA

What is the Equivalent DFA ?

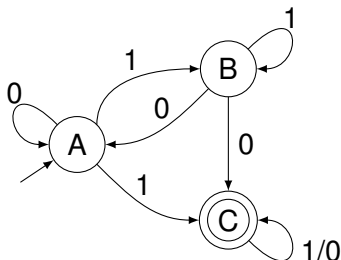


state ↓ → input characters

	0	1
A	A	BC
B	AC	B
C	C	C

# Example NFA to DFA

What is the Equivalent DFA ?

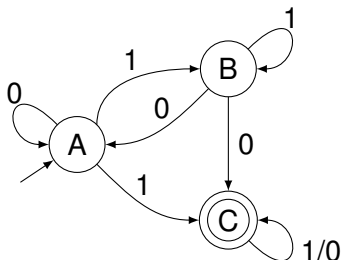


state ↓ → input characters

	0	1
A	A	BC
B	AC	B
C	C	C
AC		
BC		

# Example NFA to DFA

What is the Equivalent DFA ?

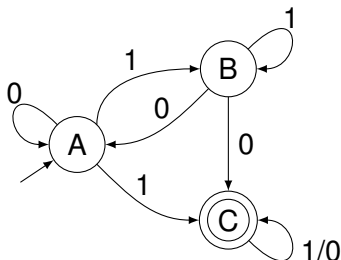


state ↓ → input characters

	0	1
A	A	BC
B	AC	B
C	C	C
AC	AC	BC
BC	AC	BC

# Example NFA to DFA

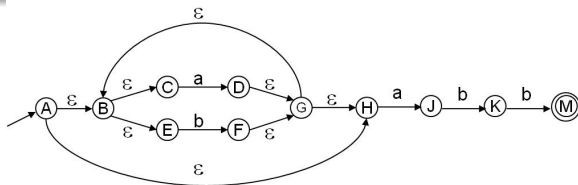
What is the Equivalent DFA ?



state ↓ → input characters


	0	1
A	A	BC
B	AC	B
C	C	C
AC	AC	BC
BC	AC	BC
AB	x	x
ABC	x	x

# Algorithm Illustrated: Converting NFA to DFA



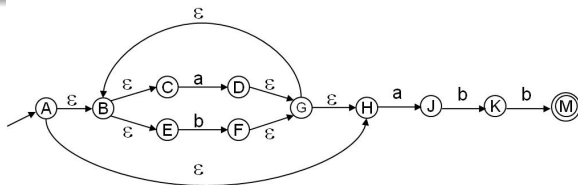
	$\epsilon$	$a$	$b$
A			
B			
C			
D			
E			
F			
G			
H			
J			
K			
M			

# Step 1: Construct the Table



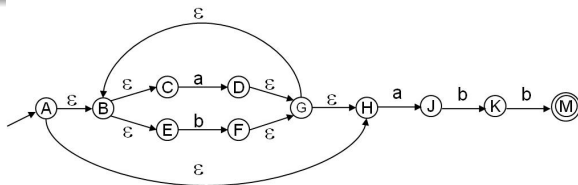
	$\epsilon$	a	b
A	BH		
B	CE		
C		D	
D	G		
E			F
F	G		
G	BH		
H		J	
J			K
K			M
M			

# Step 2: Construct $\epsilon$ -closure



	$\epsilon$	a	b
A	BHCE		
B	CE		
C		D	
D	GBHCE		
E			F
F	GBHCE		
G	BHCE		
H		J	
J			K
K			M
M			

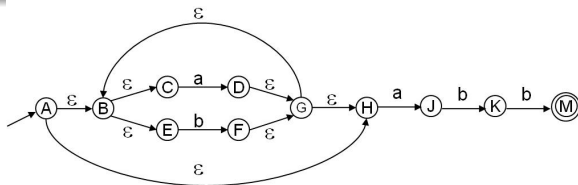
# Step 3: Update Other Columns



	$\epsilon$	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

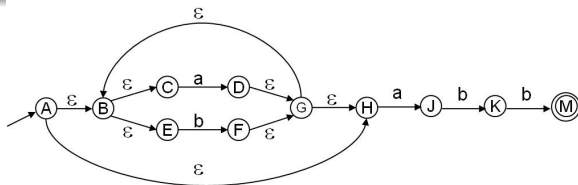


# Step 4: Construct a New Table



	$\epsilon$	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

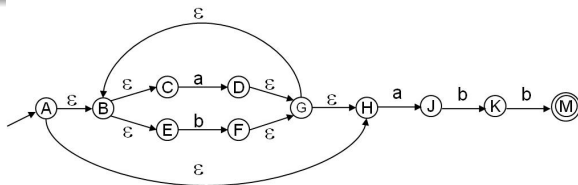
# Step 4: Construct a New Table



	$\epsilon$	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

	a	b
ABHCE	DJ	F

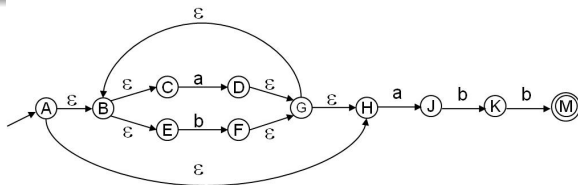
# Step 4: Construct a New Table



	$\epsilon$	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F


# Step 4: Construct a New Table



	$\epsilon$	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM

# Step 4: Construct a New Table



	$\epsilon$	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			


	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM
FM	DJ	F


# Step 5: Generate the DFA

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM
FM	DJ	F

# Step 5: Generate the DFA

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM
FM	DJ	F

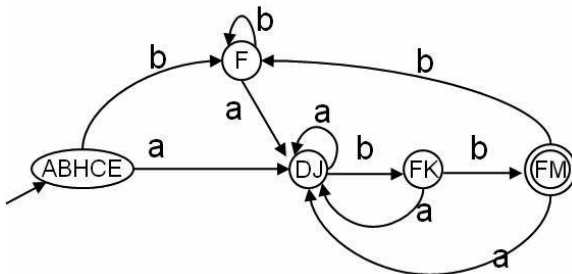
 Note: the number of states is not minimized



# Step 5: Generate the DFA

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM
FM	DJ	F

□ Note: the number of states is not minimized



□ States ABHCE and F can be merged! Why?



# NFA to DFA. Space Complexity

- ❑ An NFA may be in many states at any time
- ❑ If NFA has  $N$  states, how many DFA states?
  - NFA must be in some subset of those  $N$  states
  - How many non-empty subsets are there ?
    - $2^N - 1$  many states
- ❑ The resulting DFA has  $O(2^N)$  space complexity
  - Remember this is big-O. Typically much less.


# NFA to DFA Time Complexity

- ❑ A DFA can be implemented by a 2D table  $T$ 
  - One dimension is “states”, the other is “input characters”
  - For  $S_a \xrightarrow{c} S_b$ , we have  $T[S_a, c] = S_b$
- ❑ DFA execution
  - If the current state is  $S_a$  and input is  $c$ , then read  $T[S_a, c]$
  - Update the current state to  $S_b$ , assuming  $T[S_a, c] = S_b$
  - Requires  $O(|X|)$  steps, where  $|X|$  is the length of input
- ❑ NFA execution
  - At a given step, there is a set of possible states, up to  $N$
  - On input  $c$ , must access table for each possible state to get next set of possible states
  - Requires  $O(|X| * N)$  steps

# Implementation in Practice

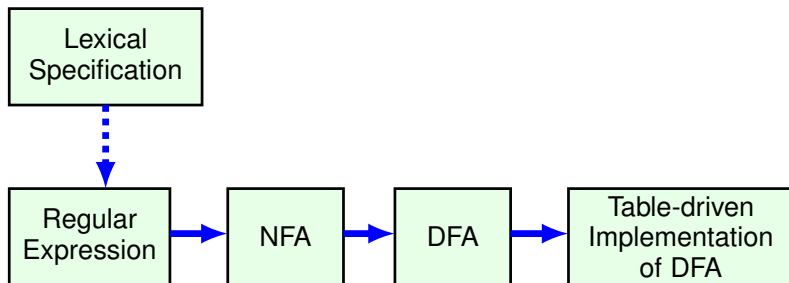
## GNU **lex**

- Convert regular expression to NFA
- Convert NFA to DFA
- Perform DFA state minimization to reduce space
- Generate transition table from DFA
- Perform table compression to further reduce space

 Most other automated lexers also trade off space for speed by choosing DFA over NFA

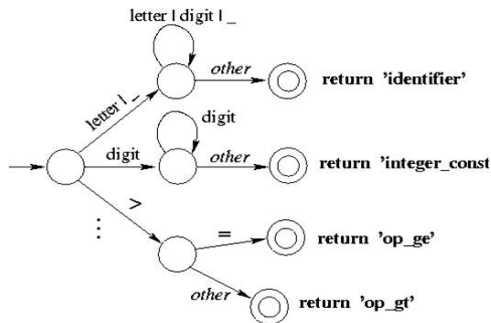
# From RE to FA

## Our implementation sketch



# Structure of a Scanner Automaton

■ A scanner recognize multiple REs



# How many characters should be matched?

- In general, find the longest match possible
- If same length, first rule in lex file takes precedence


Example:

on input **123.45**, we match it as

(numConst, 123.45)

rather than

(numConst, 123), (dot, "."), (numConst, 45)



# Beyond Regular Languages

- ❑ Regular languages are expressive enough for tokens
- ❑ Can express identifiers, strings, comments, etc.
- ❑ However, it is the least expressive formal language .
  - Many languages are not regular
    - C programming language is not
    - XML or JSON is not
    - “(((...)))” is also not
  - Finite automata cannot express nested structures
    - Can only remember a finite number of “(”s encountered
- ❑ We need a more powerful language for parsing
  - In the next lecture, we will discuss **context-free languages**